



THE UNIVERSITY OF CHICAGO

BALANCE SHEET CAPACITY AND CROSS-STRATEGY
TREASURY ARBITRAGE: SLR EXCLUSION EVENT
STUDY

By
Bailey Meche

April 2026

A paper submitted in partial fulfillment of the requirements for
the Master of Arts degree in the Master of Arts in
Computational Social Science

Faculty Advisor: Jeremiah Bejarano

Preceptor: Melissa Tartari

Abstract

Regulatory leverage constraints are a leading candidate explanation for limited intermediation in U.S. Treasury markets. In this paper, I exploit the Federal Reserve’s temporary exclusion of U.S. Treasury securities and Federal Reserve deposits from the Supplementary Leverage Ratio (SLR) denominator (April 1, 2020 to March 31, 2021) to test whether SLR relief differentially compressed Treasury-based arbitrage spreads relative to non-Treasury controls. Using a pooled difference-in-differences design across 17 daily arbitrage spread series in four strategy classes (Treasury spot-futures, TIPS-Treasury, covered interest parity, and equity spot-futures), I find that during the relief window Treasury spreads compressed by 10.76 basis points more than non-Treasury spreads relative to the 2019 pre-COVID baseline, and re-widened at expiry by 3.84 basis points in the subsequent 60 trading days. Non-Treasury series show no comparable break at the SLR event dates. This behavior is consistent with balance-sheet segmentation rather than a general intermediary funding shock and yields an arbitrage-spread-based estimate of the leverage ratio friction in Treasury markets.

Keywords: Supplementary Leverage Ratio; Treasury market arbitrage; Segmented arbitrage; Intermediary balance sheets; Difference-in-differences.

1 Background

The Supplementary Leverage Ratio is a non-risk-weighted capital requirement on large U.S. bank holding companies, requiring Tier 1 capital to exceed a minimum fraction of total leverage exposure. Because the denominator is unweighted, U.S. Treasuries and deposits at the Federal Reserve consume the same balance sheet capacity as corporate loans or risky securities: for an intermediary whose SLR binds, an additional dollar of Treasuries displaces a dollar of alternative balance sheet use. Under frictionless intermediation the arbitrage spreads studied here would be zero; in the segmented-arbitrage framework of Siriwardane et al., 2025, their magnitudes price the segment-specific balance sheet capacity required to intermediate the trade, since arbitrageurs must simultaneously hold the cheap leg and short the expensive leg, and that inventory consumes leverage.

The Federal Reserve’s 2020–2021 SLR exclusion removed this leverage cost specifically for Treasury and reserve holdings at covered bank holding companies, effective April 1, 2020 and expiring March 31, 2021, while leaving the capacity cost for foreign-currency, equity, and other positions unchanged. This regulatory asymmetry generates this paper’s treatment-control structure directly from the segmentation framework: the Treasury spot-

futures and TIPS–Treasury series are within the regulation’s scope and intermediated by segments whose balance sheet cost fell (treatment), while the CIP and equity spot-futures series are outside the regulation’s scope and intermediated by segments whose balance sheet cost did not (control). The SLR window therefore delivers a cross-strategy experiment in which the treated spreads’ balance sheet cost was reduced but the controls’ was not, yielding a testable differential prediction without requiring the broader balance sheet environment to be held constant.

Hypothesis

The SLR exclusion generates two testable predictions concerning differential Treasury price impact: spreads re-widen at SLR expiry, and spreads compress during the relief period relative to the pre-COVID baseline.

The primary leg of this hypothesis is that the spreads of Treasury arbitrage strategies widen at expiry of the SLR exclusion period. On its expiration of March 31, 2021, covered banks faced restored leverage costs on their Treasury and reserve holdings. If these costs constrain the balance sheet-intensive spot-futures basis and TIPS-inflation swap trades, then Treasury-based arbitrage spreads should widen relative to non-Treasury spreads in the post-expiry window. CIP deviations and equity spot-futures are not directly subject to the SLR exclusion: CIP basis trades do consume dealer balance sheet and may receive indirect spillovers from freed capacity (which would bias the Treasury differential toward zero), but the exemption explicitly targeted Treasury and reserve holdings rather than cross-currency positions or equity derivatives. Equity spreads are mechanically linked to short-term Treasury yields (see footnote in Section 3.1), which introduces a potential conservative attenuation. Non-Treasury spreads should show no comparable break at the SLR event dates. This is the paper’s primary source of evidence because the 60-trading-day pre-period for the exit event is free of the acute market dislocations that confound the entry window. Note, however, that this pre-period falls entirely within the SLR relief itself; the exit evidence therefore rests on the cross-strategy comparison at the expiry date rather than a clean pre-treatment baseline. The 2019 pre-COVID clean DiD provides the untreated baseline evidence.

The second leg to this hypothesis is Treasury arbitrage strategies compressing during the SLR exclusion period. During this period (April 2020 - March 2021), Treasury securities were cost-free from the perspective of the SLR denominator. If balance sheet capacity is the binding constraint on Treasury arbitrage, Treasury-based spreads should compress more than non-Treasury spreads relative to the pre-COVID 2019 baseline. This is tested using a full-period difference-in-differences that excludes the February–March 2020 COVID crash from the estimation sample. Because the entry event window coincides with the acute phase

of the COVID crisis—during which the Federal Reserve simultaneously cut rates to zero, launched emergency repo operations, and expanded bilateral swap lines—the event-window estimate for entry is likely contaminated by concurrent interventions and is reported with an explicit caveat. The full-period specification is the preferred design for measuring the relief effect.

2 Literature Review

The arbitrage spreads studied in this paper, defined as “riskless returns in excess of riskless rate” (Siriwardane et al., 2025), are persistent, economically large, and have been the subject of a significant research agenda in financial economics. This section traces that literature chronologically to motivate two claims: first, that the SLR exclusion should differentially compress Treasury-based spreads during the 2020–2021 relief window; and second, that the non-Treasury spreads in the panel serve as a natural falsification benchmark against general intermediary-health shocks.

The modern limits-of-arbitrage literature begins with Shleifer and Vishny, 1997, who show that when arbitrage is delegated to specialized professionals managing outside capital, performance-based capital flows withdraw resources precisely when mispricings have widened and expected returns are highest, leaving anomalies to persist. Gromb and Vayanos, 2002 formalize this intuition in a general equilibrium with financially constrained arbitrageurs who must collateralize each leg of a basis trade separately so that wealth shocks amplify into wider spreads. Brunnermeier and Pedersen, 2009 embed this logic in a dynamic setting in which an asset’s market liquidity and the funding liquidity of its marginal arbitrageur reinforce each other through endogenous margins, producing destabilizing margin and loss spirals. Gârleanu and Pedersen, 2011 close the pre-crisis phase by deriving the pricing equation that essentially summarizes the entire subsequent empirical literature: the basis between two securities with identical cash flows equals the product of their margin differential and the shadow cost of arbitrageur capital. The arbitrage spreads studied here admit a similar decomposition.

After 2008, a successor literature moved intermediary balance sheet health from the margin to the center of asset pricing. He and Krishnamurthy, 2013 construct a dynamic equilibrium in which the marginal investor in risky assets is a capital-constrained specialist, so that risk premia spike non-linearly when the specialist’s equity constraint binds. Adrian et al., 2014 provide complementary empirical support, showing that shocks to broker-dealer leverage price the cross-section of stock and Treasury returns in a single-factor stochastic discount factor. He et al., 2017 extend this framework to seven sophisticated asset classes using the equity capital ratio of the New York Fed primary dealers and document a consis-

tent positive price of risk of comparable magnitude across markets. The organizing premise across this literature is that a single intermediary kernel prices intermediated risky assets.

An empirical literature documented persistent, economically large law-of-one-price violations in the precise markets this paper studies. Fleckenstein et al., 2014 show that inflation-swap-stripped TIPS trade systematically cheap to matched nominal Treasuries, with aggregate mispricing exceeding \$56 billion and individual deviations above \$20 per \$100 face value. The authors point to this persistence as incompatible with frictionless arbitrage. Du et al., 2018 extend the puzzle to global currency markets, documenting persistent post-2008 CIP violations across all G-10 pairs and identifying a quarter-end signature that points to bank leverage regulation as the proximate cause. Andersen et al., 2019 supply a corporate-finance microfoundation: the funding value adjustments that dealer banks subtract from derivatives mark-to-market are debt-overhang costs borne by shareholders when new positions are financed with unsecured debt, rationalizing a structural wedge between competitive mid-market and dealer-quoted prices. Fleckenstein and Longstaff, 2020 unify these findings under a common label “the cost of renting intermediary balance sheet space” and show that an analogous funding basis exists in Treasury futures and a broad range of cash-derivative pairs, with magnitudes that predate and transcend any single regulatory regime.

The synthesis of this three-decade arc, and the theoretical backbone of this paper, is Siriwardane et al., 2025. Studying daily pricing of 32 arbitrage spreads across seven strategy classes over 2010–2020, the authors report an average pairwise correlation of only 22% across spreads, only 5% across spread changes, and eleven principal components required to explain 90% of cross-sectional variation. Such low correlations are “inconsistent with canonical intermediary asset pricing models” of the representative-intermediary type: under a single intermediary stochastic discount factor, all arbitrages should load on a common factor, and the data reject this. The authors’ interpretation is that arbitrage is segmented along two dimensions: balance-sheet segmentation and funding segmentation. Balance-sheet segmentation means that different regulatory and internal-risk constraints bind on the balance sheets that intermediate different trades: primary dealers, fixed-income hedge funds, and equity-derivatives desks each face a distinct binding constraint, so a shock to any one of them propagates only to the trades that type intermediates. Funding segmentation means that different trades depend on different funding sources: repo for secured Treasury-based trades, unsecured dealer funding for CIP and equity-based trades. Rather than a single intermediary kernel, arbitrage spreads are priced by a constellation of segment-specific shadow costs of capital. Different hedge funds matter for different Treasury basis trades; broker-dealer balance sheets dominate equity spot-futures; global bank funding conditions intermediate CIP.

A corollary of segmentation, exploited by Siriwardane et al., 2025 in their own 2016 money market fund reform event study, is that a regulatory shock affecting a particular funding source or class of intermediary should propagate to the spreads that segment intermediates, and not to others. The authors frame their MMF analysis as “formally closer to a placebo test” than a difference-in-differences: under segmentation there is no reason to expect parallel trends across unrelated segments, so the absence of a spread response in unaffected segments is itself the identifying evidence. This paper applies the same cross-strategy identification logic to the 2020–2021 SLR exclusion.

3 Data and Methods

3.1 Data

The following daily panels of arbitrage spreads for 2019–2021 are employed, with 2019 serving as the pre-COVID baseline and 2020–2021 as the treatment and post-treatment windows. Figure 1 plots the time-series evolution of all four strategy classes across the full sample and Table 1 summarizes all data sources. All series are measured in basis points. Controls (merged by date) include: secured funding rates and wedges (SOFR, TGCR–SOFR), broad risk indicators (VIX, HY OAS, Baa–10y), and Treasury issuance volumes by tenor.

TIPS–Treasury Arbitrage. As in Fleckenstein et al., 2014 and Siriwardane et al., 2025, the spread $W_{t,\tau} = (y_{t,\tau}^{\text{TIPS}} + \pi_{t,\tau}^{\text{swap}}) - y_{t,\tau}^{\text{nom}}$ for $\tau \in \{2, 5, 10\}$ years. Here $y_{t,\tau}^{\text{TIPS}} + \pi_{t,\tau}^{\text{swap}}$ is the synthetic nominal yield.

Treasury Spot–Futures Arbitrage. For each of the 2-, 5-, and 10-year Treasuries, I use the first-deferred futures contract to compute a no-arbitrage implied yield (spot-futures parity) and subtract the matching OIS rate. I retain the 2-, 5-, and 10-year tenors, for which data coverage and liquidity are most reliable across the 2019–2021 sample. The construction follows Siriwardane et al., 2025, who apply the same methodology for the 2-, 5-, 10-, 20-, and 30-year tenors over 2010–2020.

CIP Arbitrage. Following Du et al., 2018, for each of eight G-10 currencies (AUD, CAD, CHF, EUR, GBP, JPY, NZD, SEK) I define the 3-month CIP basis as the difference between the dollar OIS rate and the synthetic forward-implied rate from currency spot and forward quotes. I examine all eight individually and present pooled and series-level results; the mean across currencies is used for strategy-class summary statistics.

Equity Spot–Futures Arbitrage. For three equity indexes (SPX, INDU, NDX) I calculate the rate implied by equity index futures (accounting for dividends) and subtract the short-term Treasury yield. This spread is based on the violation of stock spot–futures parity.¹

3.2 Methods

The Federal Reserve’s temporary exclusion of U.S. Treasury securities and Federal Reserve deposits from the Supplementary Leverage Ratio (SLR) denominator provides a temporary regulatory change in the balance sheet capacity of covered bank holding companies. I use this event to test whether regulatory leverage constraints are the operative friction behind Treasury-based arbitrage dislocations. The SLR exclusion was itself an endogenous policy response to COVID-19 market stress, so the design does not rely on the exclusion being an exogenous shock to the wider economy; rather, in the spirit of the MMF reform analysis in Siriwardane et al., 2025, Section 4.2, it exploits a dated regulatory event with a known mechanism and asks whether the pattern of spread responses across strategy classes matches the theoretical prediction, with the non-Treasury series serving as a falsification benchmark. The identification of this study rests on a single maintained hypothesis: if balance sheet

¹The equity spot-futures spread mechanically depends on short-term Treasury yields through its construction. If the SLR exclusion affected Treasury yield levels, it could in principle compress the equity spread through the Treasury-yield component, biasing $\hat{\beta}$ toward zero (a conservative attenuation). However, the near-zero exit responses for equity series cannot fully resolve this concern: contamination would show up precisely as small differential effects (equity spreads moving in the same direction as Treasury spreads due to the shared yield component), which is observationally equivalent to a true null. The direction of bias depends on whether Treasury yield levels (not spreads) were affected—a distinction the current design cannot cleanly make.

Table 1: Data Sources and Coverage

Series	Construction	Source	<i>N</i> obs
<i>Panel A: Treasury-based arbitrage spreads</i>			
TIPS–Treasury (2y, 5y, 10y)	$y_{t,\tau}^{\text{TIPS}} + \pi_{t,\tau}^{\text{swap}} - y_{t,\tau}^{\text{nom}}$	Bloomberg / Fed H.15	3,750/tenor
UST spot-fut. (2y, 5y, 10y)	Futures-implied yield minus OIS	Bloomberg	4,967/tenor
<i>Panel B: Non-Treasury arbitrage spreads</i>			
CIP 3m (8 G-10 currencies)	Dollar OIS minus synthetic forward rate	Bloomberg	3,910/tenor
Equity SF (SPX, NDX, INDU)	Futures-implied forward minus T-bill	Bloomberg	3,730/tenor
<i>Panel C: Controls</i>			
VIX, HY OAS, Baa–10y	Daily index levels	CBOE / ICE BofA / FRED	794
SOFR, TGCR–SOFR	Overnight secured rates	NY Fed STFM	784
Treasury issuance	Rolling 7/14/30-day total	FiscalData Treasury	794

Notes. Estimation sample: January 1, 2019 to December 31, 2021. The final stacked estimation panel contains 68,620 series-day observations across 17 series (after merging to controls over 2019–2021).

constraints drive segmentation in Treasury markets specifically, then the onset and expiry of a relief that removes Treasuries—but not other assets—from the leverage denominator should compress Treasury-based spreads relative to non-Treasury spreads, and then reverse that compression when the relief ends.

The primary specification is a pooled difference-in-differences regression estimated on a cross-sectional panel of 17 arbitrage spread series spanning four strategy classes: U.S. Treasury spot-futures (three tenors), TIPS-Treasury (three tenors), covered interest parity deviations for eight G-10 currencies, and equity spot-futures for three indices. I define the dependent variable as $|W_{i,t}|$, the absolute magnitude of the arbitrage spread for series i on trading day t , measured in basis points. Working with the absolute value follows Siriwardane et al., 2025, who note that the sign of an arbitrage spread depends on whether




Figure 1: Time-series of arbitrage spreads, 2019–2021. Each panel shows the cross-series mean (solid line) and ± 1 standard deviation band (shaded) of absolute spread $|W|$ in basis points, smoothed with a 5-day rolling average. The gray shaded region denotes the SLR exclusion relief period (April 1, 2020 – March 31, 2021). Dashed vertical lines mark the entry (2020-04-01) and exit (2021-03-31) dates.

intermediaries are net long or short a given leg of the trade, so the magnitude is the natural measure of mispricing.

The classification of series into treatment and control groups is determined by the mechanism of the SLR relief. The interim final rule excluded U.S. Treasury securities and deposits at Federal Reserve Banks from the SLR denominator for covered bank holding companies, with no analogous relief for foreign-currency-denominated assets, equity derivatives, or other securities. I therefore classify the Treasury spot-futures and TIPS-Treasury series as *Treasury-based* (treated, $TreasuryBased_i = 1$) and the CIP and equity spot-futures series as *non-Treasury* (control, $TreasuryBased_i = 0$). This classification is primarily determined by the regulation’s direct scope: the SLR rule excluded Treasuries and reserves, not FX derivatives or equity futures, from the denominator. General equilibrium spillovers—such as freed balance sheet capacity being reallocated toward CIP trades—cannot be ruled out in principle, but would bias $\hat{\beta}$ toward zero, making the estimates conservative.




Figure 2: UST spot-futures arbitrage spread by tenor, 2019–2021. Lines show the absolute spread $|W|$ for 2-year (solid), 5-year (dashed), and 10-year (dotted) maturities, smoothed with a 5-day rolling average. Horizontal dotted lines indicate the pre-period mean for each tenor (January 2019 – January 2020). The gray shaded region and dashed vertical lines are as in Figure 1.

Let Post_t be an indicator equal to one for trading days in the post-event window, and let α_i denote a full set of series fixed effects. My baseline specification is:

$$|W_{i,t}| = \alpha_i + \gamma \text{Post}_t + \beta(\text{Post}_t \times \text{TreasuryBased}_i) + \delta' \mathbf{X}_t + \varepsilon_{i,t}, \quad (1)$$

where $\mathbf{X}_t$ is a vector of daily controls including the CBOE Volatility Index (VIX), the ICE BofA High-Yield OAS, and the Baa-to-10-year Treasury credit spread (TOTAL specification). The DIRECT specification augments this with the SOFR rate level, the TGCR-to-

Table 2: Summary Statistics for Arbitrage Spreads

Strategy	Series	N	Mean $\bar{W}$	Median W	Mean $ W $	$Q_{0.05}(W)$	$Q_{0.95}(W)$
<i>Panel A: Treasury-based arbitrages ($\text{TreasuryBased} = 1$)</i>							
TIPS–Treasury	2y, pre	312	16.6	19.5	19.4	−9.7	36.8
	2y, relief	250	16.8	22.3	21.4	−16.7	33.7
	2y, post	190	11.7	13.4	14.8	−11.1	27.8
	5y, pre	312	16.3	16.8	16.3	6.4	24.3
	5y, relief	250	13.8	16.2	14.1	−0.4	22.3
	5y, post	190	8.7	9.0	9.0	−0.4	16.9
	10y, pre	312	24.1	23.5	24.1	18.8	29.2
	10y, relief	250	22.9	26.3	22.9	1.1	31.4
	10y, post	190	10.3	11.4	10.4	1.1	19.2
UST spot–fut.	2y, pre	314	−24.4	−24.3	25.6	−38.0	−9.6
	2y, relief	251	−5.8	−6.0	7.6	−16.1	9.4
	2y, post	192	10.0	8.0	12.6	−3.2	27.4
	5y, pre	314	−25.4	−26.3	26.4	−39.9	−6.6
	5y, relief	251	−6.3	−6.5	7.7	−17.8	5.9
	5y, post	192	5.5	5.3	5.8	−0.4	12.8
	10y, pre	314	−20.8	−23.4	24.7	−39.1	6.9
	10y, relief	251	−5.4	−6.1	8.2	−16.0	7.7
	10y, post	192	10.1	8.9	10.2	2.8	20.1
<i>Panel B: Non-Treasury arbitrages ($\text{TreasuryBased} = 0$, means across tenors/currencies)</i>							
CIP 3m (8 curr.)	pre	326 avg.	—	—	30.3	—	—
	relief	261 avg.	—	—	25.9	—	—
	post	197 avg.	—	—	17.2	—	—
Equity SF (3 idx.)	pre	310 avg.	—	—	32.6	—	—
	relief	248 avg.	—	—	45.3	—	—
	post	189 avg.	—	—	50.9	—	—

Notes. Spread statistics are in basis points. $\bar{W}$ denotes the mean signed spread; $|W|$ denotes the mean absolute spread. Regime boundaries: pre = January 1, 2019 to March 31, 2020; relief = April 1, 2020 to March 31, 2021; post = April 1 to December 31, 2021. CIP and equity rows report cross-series averages within each regime. All series are winsorized at the within-series first and 99th percentiles prior to analysis; see Table 9. The ‘pre’ regime includes February–March 2020 (the acute COVID crash), which are excluded from the DiD regressions in Table 5 Panel B. Pre-period means for UST spot-futures are therefore modestly inflated by COVID-period observations relative to the 2019-only baseline used in the DiD.

SOFR spread, and rolling 7-, 14-, and 30-day Treasury issuance volumes (in billions, lagged one day for point-in-time discipline). Unless otherwise noted, the DIRECT specification includes all five additional controls throughout. Series fixed effects α_i absorb all time-invariant differences in average spread levels across instruments. β is identified from the differential within-series variation across the two groups (Treasury-based vs. non-Treasury): it measures how much more Treasury-based spreads changed around the event relative to non-Treasury spreads, after controlling for level differences and common daily shocks. I do not include date fixed effects because date indicators would absorb the Post_t variation that identifies γ and β ; instead, this study controls for common daily shocks through the vector $\mathbf{X}_t$ of observable risk factors. The coefficient γ captures the average post-event change in




Figure 3: Distribution of absolute arbitrage spreads $|W|$ by regime. Each box shows the interquartile range with median line; whiskers extend to $1.5 \times \text{IQR}$. Treasury-based series (UST spot-futures and TIPS-Treasury) are shown in the left panels; non-Treasury series (CIP deviations and equity spot-futures) in the right panels. Regimes: Pre = January 2019 – March 2020; Relief = April 2020 – March 2021; Post = April 2021 – December 2021 (February–March 2020 excluded from DiD regressions but included here for visual completeness).

non-Treasury spread levels, while β captures the incremental post-event change in Treasury-based spreads relative to non-Treasury spreads. Standard errors are computed using the Newey-West heteroskedasticity and autocorrelation consistent (HAC) estimator with five lags, which accommodates the high serial persistence documented in Table 2 (daily AR(1) coefficients exceeding 0.93 for most series).² I report Driscoll and Kraay, 1998 standard errors as a robustness check in Section 4.5.

I estimate (1) separately for the entry event (April 1, 2020) and the exit event (March 31, 2021), and separately for symmetric windows of ± 20 and ± 60 trading days around each event date. As I discuss below, the exit event provides sharper evidence because the pre-period of the exit window—falling entirely within the post-COVID normalization period of late 2020—satisfies the parallel-trends assumption in a way the entry window cannot.

Next, a full-period difference-in-differences specification is estimated using the 2019 calendar year as the pre-COVID baseline, excluding the February-March 2020 COVID crash from the sample. In this specification, Relief_t equals one during April 1, 2020 through March 31, 2021, PostRelief_t equals one from April 1, 2021 onward, and the model includes interactions of both with TreasuryBased_i :

$$\begin{aligned} |W_{i,t}| = & \alpha_i + \mu_1 \text{Relief}_t + \mu_2 (\text{Relief}_t \times \text{TreasuryBased}_i) \\ & + \mu_3 \text{PostRelief}_t + \mu_4 (\text{PostRelief}_t \times \text{TreasuryBased}_i) + \boldsymbol{\delta}' \mathbf{X}_t + \varepsilon_{i,t}. \end{aligned} \quad (2)$$

The coefficients μ_2 and μ_4 capture the differential Treasury-based spread response during and after the relief period, respectively, relative to the 2019 pre-period baseline.

4 Results

The difference-in-differences design requires that, absent the SLR event, Treasury-based and non-Treasury spreads would have followed parallel trends. I test this directly within the 2019 pre-treatment year by interacting TreasuryBased_i with quarterly indicators Q_t and estimating

$$|W_{i,t}| = \alpha_i + \sum_{q=2}^4 \lambda_q Q_{q,t} + \sum_{q=2}^4 \delta_q (Q_{q,t} \times \text{TreasuryBased}_i) + \boldsymbol{\gamma}' \mathbf{X}_t + \varepsilon_{i,t}, \quad (3)$$

on the 2019 calendar year. Results are shown in Table 3. A joint Wald test of $H_0: \delta_2 = \delta_3 = \delta_4 = 0$ yields $F(3, 4,306) = 17.04$, $p < 0.001$, decisively rejecting the null. The Treasury-

²The one exception is UST SF 10y, with a daily AR(1) coefficient of 0.881. The remaining 16 of 17 series have AR(1) coefficients above 0.93; all 17 exceed 0.88.

based and non-Treasury spread groups had significantly different seasonal dynamics in 2019. This is an important limitation of the parallel-trends assumption for the DiD design. The magnitude of the treatment effect ($\hat{\mu}_2 = -10.76$ bps) substantially exceeds these seasonal fluctuations, but I interpret the DiD estimate as lower-bound evidence rather than a clean identification of the causal effect. The secular co-movement of both spread classes over the post-GFC decade Siriwardane et al., 2025 provides qualitative support for the claim that structural trends were similar across classes before the SLR event, even if the unconditional seasonal test rejects.

The $W = 60$ exit event window cannot provide an additional parallel-trends test, because the 60-trading-day pre-period falls entirely within the SLR relief period itself. With this in mind, I estimate a dynamic version of (1) replacing Post_t with event-time bin indicators to characterize the within-window spread dynamics around the exit date. The results in

Table 3: Quarterly Pre-Trend Test: TreasuryBased $\times$ Quarter Interactions (2019)

Interaction	$\hat{\delta}_q$ (bps)	t -stat	
$Q_2 \times \text{TreasuryBased}$	-0.44	-0.29	
$Q_3 \times \text{TreasuryBased}$	+6.25	+3.77	***
$Q_4 \times \text{TreasuryBased}$	-5.58	-3.00	***
<i>Joint Wald test: $F(3, 4,306) = 17.04, p < 0.001$</i>			

Notes. Q_1 (January–March 2019) is the omitted reference quarter. Each $\hat{\delta}_q$ measures the differential level of Treasury-based versus non-Treasury spreads in quarter q relative to Q_1 , conditional on series fixed effects and TOTAL controls (VIX, HY OAS, Baa-10y). $N = 4,332$ series-day observations; HAC SE (Newey-West, 5 lags). *** $p < 0.01$.

Table 4: Event-Study Dynamics Around SLR Exit, $W = 60$

Event-time bin	$\hat{\beta}_k$ (bps)	SE	t
<i>Pre-event bins (within SLR relief period; see text)</i>			
$[-60, -51]$	-1.81	1.72	-1.05
$[-50, -41]$	-4.72***	1.67	-2.83
$[-40, -31]$	-4.37***	0.97	-4.49
$[-30, -21]$	-2.90**	1.14	-2.54
$[-20, -11]$	-3.91***	1.38	-2.83
$[-10, -1]$	-3.48***	0.91	-3.82
F -test (6 pre-event leads)	$F = 9.92, p < 0.001$		
<i>Post-event lags</i>			
$[+0, +9]$	+0.56	1.06	+0.53
$[+10, +19]$	+0.69	0.88	+0.78
$[+20, +29]$	-0.70	0.91	-0.77
$[+30, +39]$	-2.77***	0.89	-3.12
$[+40, +49]$	-1.17	0.80	-1.46
$[+50, +60]$	+2.19**	0.90	+2.43

Notes. HAC standard errors (Newey-West, 5 lags). *** $p < 0.01$; ** $p < 0.05$.

Table 4 show significant negative pre-event interaction coefficients were actively compressing relative to their all-period mean throughout the pre-window: this is the treatment effect in progress, not a pre-trend divergence. This non-monotone pattern for post-event lag coefficients is consistent with within-window mean-reversion dynamics at the series fixed-effect level and does not contradict the pooled level-shift estimate of $\hat{\beta} = +3.84$ bps, which directly compares the pre-expiry and post-expiry within-window means.

Table 5 reports estimates of (1) for the exit and entry event, across both window sizes and both control specifications. I begin with the exit event. The exit specification uses a pre-event window that falls between October 2020 and March 2021. This is a period of sustained post-COVID market normalization, well after the acute dislocations of March-April 2020. Figure 4 plots each series' mean absolute spread demeaned by its pre-event mean (days -60 to -1), so that the cross-strategy divergence at event time zero is immediately visible: Treasury-based spreads re-widen sharply while non-Treasury spreads continue their gradual compression, isolating $\hat{\beta}$ as the differential response at expiry.

The estimated interaction coefficient at $W = 60$ under the TOTAL specification is

Table 5: SLR Exclusion Event Study: Pooled Jump Regressions

	Entry event (2020-04-01) [†]				Exit event (2021-03-31)			
	$W = 20$		$W = 60$		$W = 20$		$W = 60$	
	TOTAL	DIRECT	TOTAL	DIRECT	TOTAL	DIRECT	TOTAL	DIRECT
<i>Panel A: Pooled regressions</i>								
Post	−0.84 (4.77)	+1.45 (7.03)	−7.08** (3.05)	−2.76 (5.54)	−2.62*** (0.77)	−2.28*** (0.68)	−3.09*** (0.83)	−3.23*** (0.85)
Post $\times$ TreasuryBased	−2.16 (6.13)	−2.23 (6.14)	−4.54 (3.00)	−4.52 (3.00)	+3.91*** (0.92)	+3.90*** (0.92)	+3.84*** (0.89)	+3.84*** (0.88)
N	697	697	2,057	2,057	697	697	2,057	2,057
R^2	0.61	0.61	0.55	0.55	0.96	0.96	0.89	0.89
<i>Panel B: Robust DiD (2019 pre-COVID baseline, excl. Feb–Mar 2020)</i>								
	Relief period		Post-relief period					
	TOTAL	DIRECT	TOTAL	DIRECT				
Non-Treasury baseline	−1.34 (1.13)	−11.03*** (2.97)	−6.92*** (1.01)	−14.87*** (2.80)				
$\times$ TreasuryBased	−10.76*** (1.15)	−10.77*** (1.14)	−9.03*** (1.16)	−9.04*** (1.16)				
N	12,326		12,326					
R^2	0.516		0.521				0.521	

Notes. HAC standard errors (Newey-West, 5 lags) in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. [†] The entry pre-period spans the COVID-19 crash peak; entry estimates are contaminated by contemporaneous Fed interventions. See Section 4.2 and Panel B (Robust DiD) for COVID-robust estimates.

$\hat{\beta} = +3.84$ basis points. The DIRECT specification, which adds SOFR and TGCR-SOFR controls, yields $\hat{\beta} = +3.84$ basis points. These estimates are identical to two decimal places; the negligible difference between the TOTAL and DIRECT specifications is informative: adding secured overnight funding rate controls does not move the Treasury differential at all, confirming that the differential widening of Treasury spreads at expiry operates through a channel orthogonal to secured funding market conditions. This stability rules out one specific alternative—that post-expiry Treasury spread widening reflects a contemporaneous tightening of secured repo markets—and is consistent with the balance sheet constraint channel as the operative mechanism, though it does not rule out other coincident factors such as anticipatory portfolio rebalancing or shifts in hedge fund leverage.

The results are equally stable at the narrower $W = 20$ window: $\hat{\beta} = +3.91$ basis points under TOTAL and $+3.90$ basis points under DIRECT. The insensitivity to window choice further supports the conclusion that the Treasury differential reflects a discrete regime shift rather than noise.

The non-Treasury baseline coefficient $\hat{\gamma}$ at $W = 60$ is -3.09 basis points and -3.23 basis points. Non-Treasury spreads were therefore continuing to compress slightly in the




Figure 4: Event-Window Spread Paths Around SLR Exit (March 31, 2021). Daily mean absolute spread, demeaned by each series’ pre-event mean (days -60 to -1), plotted from -60 to $+60$ trading days around the SLR exit date. Left panel: Treasury-based series (UST spot-futures and TIPS-Treasury). Right panel: non-Treasury series (CIP deviations and equity spot-futures). Treasury-based spreads re-widen sharply at event time zero while non-Treasury spreads continue a gradual compression, consistent with balance sheet segmentation.

post-expiry window, reflecting the ongoing secular normalization of cross-currency bases documented by Du et al., 2018. This compression makes the +3.84 basis point Treasury differential even more economically striking: while CIP and equity bases continued tightening, Treasury-based spreads widened in the opposite direction. On an absolute basis, the total post-expiry movement for Treasury series is approximately $\hat{\gamma} + \hat{\beta} = -3.09 + 3.84 = +0.75$ basis points, while non-Treasury series moved -3.09 basis points. The divergence of 3.84 basis points is economically meaningful: it represents approximately 49% of the UST spot-futures average spread during the relief window (7.8 bps), suggesting that the first 60 trading days after expiry reversed nearly half of the per-period compression for the most directly affected series.³

4.1 Robust Difference-in-Differences

The event-window specification conditions on a narrow local window around each event date, which can conflate policy effects with contemporaneous shocks. This study complements the event-window estimates with the full-period specification in (2), using the 2019 calendar year as the pre-COVID baseline and excluding February and March 2020 from the sample to avoid contamination from the acute COVID crash. This yields a balanced panel across three regimes: pre, relief, and post-relief.

Table 6 reports estimates of (2) under both control specifications. The economic magnitude of $\hat{\mu}_2 = -10.76$ basis points is striking. During the relief window, Treasury-based spreads compressed by approximately 10.8 basis points more than non-Treasury spreads, relative to their common 2019 baseline. The total Treasury compression during relief is $\hat{\mu}_1 + \hat{\mu}_2 = -12.1$ basis points. For UST spot-futures—the most directly affected series—this is consistent with the raw data: average $|W|$ fell from 25.6 basis points in the pre-period to 7.8 basis points during the relief window, a 70% compression.

For the post-relief period, the interaction is $\hat{\mu}_4 = -9.03$ basis points (SE = 1.16, $t = -7.79$). This coefficient measures the Treasury differential in the post-relief period relative to 2019, and its magnitude indicates that Treasury-based spreads did not fully return to their pre-SLR-exclusion levels following expiry. The combined evidence is consistent with a partial reversion: the exit-window estimate of +3.84 basis points captures the discrete re-widening at expiry, while the full-period post-interaction of -9.03 basis points confirms that some compression persisted through year-end 2021. One interpretation is that the SLR relief, combined with the ongoing expansion of the Fed’s balance sheet through 2021, produced structural changes in Treasury market intermediation that outlasted the formal regulatory change.

³The figure is best read as an order-of-magnitude benchmark for the UST spot-futures channel specifically, not as a precise reversal percentage for the pooled DiD.

Testing whether the compression during the relief period equals the residual compression post-relief, I compute $\hat{\mu}_2 - \hat{\mu}_4 = -10.76 - (-9.03) = -1.73$ basis points. A Wald test of $H_0: \mu_2 = \mu_4$ yields $F(1, 12,301) = 2.56$, $p = 0.11$. We can firmly claim that, first, treasury-based spreads compressed relative to non-Treasury spreads during the relief period and, second, Treasury-based spreads remained substantially more compressed than their 2019 baseline through year-end 2021. The data do not establish that the SLR exclusion is the primary driver of the persistent compression, only that Treasury spreads did not return to their pre-exclusion levels within the 2021 post-relief window. Whether this persistence reflects slow-moving arbitrage capital, concurrent QE balance sheet expansion, or structural changes in Treasury market intermediation cannot be resolved within the current design.⁴ The sub-basis-point differences between TOTAL and DIRECT coefficients replicate the stability documented in the exit event-window regressions and confirm that secured funding rate controls are irrelevant to the Treasury differential in either regime.

The non-Treasury baseline does shift substantially between specifications: $\hat{\mu}_1$ moves from -1.34 bps (TOTAL) to -11.03 bps (DIRECT), a change of nearly 10 basis points.

⁴The Federal Reserve launched its Standing Repo Facility (SRF) on July 28, 2021, nearly four months after SLR expiry. The SRF's effect on spread persistence is speculative in the absence of a structural break test around July 2021, which is beyond the scope of the current paper.

Table 6: Robust Difference-in-Differences: Full-Period Estimates

	TOTAL	DIRECT
<i>Panel A: Relief period (2020-04-01 to 2021-03-30)</i>		
$\hat{\mu}_1$: Relief (non-Treasury baseline)	-1.34 (1.13)	-11.03*** (2.97)
$\hat{\mu}_2$: Relief $\times$ TreasuryBased	-10.76*** (1.15)	-10.77*** (1.14)
<i>Panel B: Post-relief period (2021-04-01 to 2021-12-31)</i>		
$\hat{\mu}_3$: Post-relief (non-Treasury baseline)	-6.92*** (1.01)	-14.87*** (2.80)
$\hat{\mu}_4$: Post-relief $\times$ TreasuryBased	-9.03*** (1.16)	-9.04*** (1.16)
N (total)	12,326	12,326
Pre-period	4,713	4,713
Relief period	4,335	4,335
Post-relief	3,286	3,286
R^2	0.516	0.521
Series FE	Yes	Yes
TOTAL controls	Yes	Yes
DIRECT controls	No	Yes

Notes. HAC standard errors (Newey-West, 5 lags) in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. See Section 4.1 for discussion.

This shift reflects the relationship between non-Treasury arbitrage spreads and the secured overnight funding rate (SOFR). Empirically, the estimated coefficient on SOFR in the non-Treasury spread regression is -0.82 bps per percentage point, meaning that higher SOFR is associated with lower non-Treasury $|W|$ across the sample. SOFR fell sharply to zero during the relief period implying that, holding all else fixed, the SOFR decline would predict non-Treasury spreads to be roughly $0.82 \times 2.4 \approx 2.0$ bps higher than in the pre-period. This back-of-envelope accounts for only part of the observed ≈ 10 bps shift in $\hat{\mu}_1$ across specifications; the remaining gap reflects the contributions of the other DIRECT controls and nonlinearities not captured by this single-coefficient calculation. A full decomposition of the baseline shift across all DIRECT controls is beyond the scope of the current paper. Importantly, this uncertainty about the non-Treasury baseline does not affect the paper’s primary claim: the Treasury interaction $\hat{\mu}_2$ is stable to within 0.01 bps across specifications, confirming that the differential compression of Treasury-based spreads is orthogonal to SOFR dynamics. This stability is the paper’s key empirical claim.

4.2 The Entry Event and COVID Contamination

The entry event (April 1, 2020) does not provide clean evidence. The $W = 60$ pre-period spans the acute phase of the COVID-19 financial crisis, during which the Federal Reserve simultaneously cut rates to zero, launched emergency repo operations, and expanded bilateral swap lines. These concurrent interventions make it impossible to isolate the SLR relief’s marginal contribution from a severely contaminated baseline. Entry event-window estimates are reported in Table 5 for completeness and flagged accordingly; they are negative in sign as expected but insignificant.

The clean DiD of Section 4.1 addresses this directly. By conditioning on the 2019 pre-COVID baseline and excluding the February–March 2020 crash period, it recovers a precisely estimated and strongly significant differential compression (see Table 6), confirming the balance sheet channel over the full relief horizon. The exit event-window and clean DiD are the paper’s two primary results; the entry event-window is supplementary.

4.3 Falsification: Non-Treasury Series as Controls

A central feature of the design is that non-Treasury series spreads serve as active falsification tests. If arbitrage is segmented by balance sheet type, a Treasury-specific regulatory shock should move Treasury-based spreads without moving non-Treasury spreads. Symmetric responses across groups would indicate a broader shock to intermediary health rather than a balance sheet channel.

The data support the falsification hypothesis at exit. Non-Treasury spreads continued

to compress slightly in the post-expiry window (Table 5), reflecting secular normalization of dollar funding markets (Du et al., 2018) and not a response to the SLR expiry. Within CIP, EUR/CHF/JPY/SEK continued to decline while CAD/NZD/GBP were small and mixed; CIP-AUD widened significantly but this is consistent with Japanese fiscal year-end repatriation flows at March 31, a calendar effect unrelated to the balance sheet channel.⁵ Equity series showed no response to the SLR expiry (Table 7, Panel D). SPX and INDU are near zero; NDX is marginally significant but in the same compression direction as other non-Treasury series, consistent with the conservative-attenuation channel described in Section 3.1 rather than any regulatory response. The SLR exclusion directly reduced the leverage cost of holding Treasuries, not equity derivatives, so the absence of an equity response is the predicted outcome under balance sheet segmentation. A joint test across equity series is recommended for future work.⁶

At entry, several CIP series compressed sharply (Table 7), while all equity series were insignificant. The CIP compressions reflect the Federal Reserve’s bilateral swap line activations as a dollar liquidity channel as opposed to the SLR relief. The DIRECT specification controls for SOFR and TGCR-SOFR precisely to separate these channels; the Treasury interaction remains stable across specifications while the non-Treasury baseline absorbs the funding-rate effect.

4.4 Series-Level Evidence

Table 7 and Figure 5 report the series-level estimates of (1) without the TreasuryBased_i interaction, using the $W = 60$ DIRECT window for the exit event. Three of six Treasury-based series show significant re-widening at expiry. The TIPS-Treasury series show heterogeneous responses. The heterogeneity across Treasury-based series is itself informative. UST spot-futures arbitrage is a high-frequency, balance sheet-intensive trade executed primarily by primary dealers through the repo market; the relief directly lowered the SLR cost of holding the long Treasury leg of this position, and the expiry restored it.

⁵Japanese institutions routinely repatriate U.S. dollars at March 31 via yen-denominated funding operations, reducing AUD-USD swap liquidity and widening the AUD basis. The elevated CIP-AUD spread persists through April–May 2021 at roughly 10–11 bps and the same pattern appears in non-SLR-exit years. Excluding CIP-AUD leaves the primary $\text{Post} \times \text{TreasuryBased}$ coefficient essentially unchanged (Table 7, note).

⁶The opposing dynamics within the non-Treasury group—CIP compressing while equity widens across regimes (Table 2; Figure 3)—mean that the pooled non-Treasury baseline $\hat{\gamma}$ averages over divergent trends. The key empirical claim is on the Treasury differential $\hat{\beta}$, not the baseline level.

TIPS-Treasury arbitrage, by contrast, relies on a more diffuse set of intermediaries including fixed-income hedge funds Siriwardane et al., 2025, Table VII and involves longer-duration positions for which balance sheet turnover is slower. The stronger and more uniform response among UST spot-futures series is consistent with the view that primary dealer balance sheets are the marginal arbitrageurs for this strategy. The TIPS-Treasury evidence is consistent with the balance sheet channel but suggests that additional intermediary heterogeneity, potentially at the hedge fund level, dampens the regulatory signal at exit.

The TIPS 2y series is the one Treasury-based observation inconsistent with the balance sheet hypothesis at expiry: its post-coefficient is -8.28 basis points, indicating continued compression rather than re-widening. Several non-mutually-exclusive explanations are consistent with this finding. The TIPS 2y spread takes positive values in this sample—the

Table 7: Series-Level Jump Estimates: Exit Event, $W = 60$, DIRECT

Series	Treas.-based	$\hat{c}_i$ (bps)	SE	t	N
<i>Panel A: UST spot-futures (treated)</i>					
UST SF 2y	Yes	+2.47	1.41	+1.76*	121
UST SF 5y	Yes	+4.05	1.36	+2.98***	121
UST SF 10y	Yes	+4.83	1.83	+2.64***	121
<i>Panel B: TIPS-Treasury (treated)</i>					
TIPS 2y	Yes	-8.28	4.12	-2.01**	121
TIPS 5y	Yes	-0.83	1.09	-0.76	121
TIPS 10y	Yes	+0.35	1.63	+0.22	121
<i>Panel C: CIP deviations (control)</i>					
CIP-AUD	No	+4.32	1.03	+4.19***	121
CIP-CAD	No	-2.69	2.14	-1.25	121
CIP-CHF	No	-4.35	1.47	-2.95***	121
CIP-EUR	No	-6.15	1.61	-3.81***	121
CIP-GBP	No	-2.92	1.84	-1.59	121
CIP-JPY	No	-4.46	1.55	-2.89***	121
CIP-NZD	No	-0.77	1.21	-0.63	121
CIP-SEK	No	-4.71	1.72	-2.74***	121
<i>Panel D: Equity spot-futures (control)</i>					
Eq. SPX	No	-0.72	1.47	-0.49	121
Eq. NDX	No	-4.47	2.64	-1.69*	121
Eq. INDU	No	-0.21	1.72	-0.12	121

Notes. Each row is a separate series-level estimate of the post-event coefficient $\hat{c}_i$ from a univariate version of equation (1) estimated on series i alone, without the TreasuryBased_i interaction—that is, $|W_{i,t}| = \alpha_i + \hat{c}_i \text{Post}_t + \delta' \mathbf{X}_t + \varepsilon_{i,t}$. Note that $\hat{c}_i$ is a simple pre/post level shift within series i , not an interaction effect; it is distinct from the pooled $\hat{\beta} = +3.84$ bps reported in Table 5. Exit event (March 31, 2021), $W = 60$, DIRECT specification. $N = 121$ trading days per series. HAC SE (Newey-West, 5 lags) *** $p < 0.01$; ** $p < 0.05$; * $p < 0.10$. Positive $\hat{c}_i$ indicates widening post-expiry; negative indicates further compression.

summary statistics in Table 2 show $\bar{W} \approx +3$ bps for TIPS 2y, so $|W| = W > 0$. A reduction in $|W|$ therefore means W itself compressed (moved toward zero from above). First,




Figure 5: Series-level post-event coefficients $\hat{c}_i$ around the SLR exit, $W = 60$, DIRECT specification. Each point is an independent per-series estimate with 95% HAC confidence interval (Newey-West, 5 lags). Treasury-based series (UST spot-futures and TIPS-Treasury, filled circles) are shown separately from non-Treasury controls (CIP and equity spot-futures, open circles). Positive values indicate spread widening post-expiry.

the Federal Reserve’s QE program was actively purchasing TIPS through 2021 at roughly \$5 billion per month, directly suppressing $y_{t,2}^{\text{TIPS}}$; since $W_{t,2} = (y_{t,2}^{\text{TIPS}} + \pi_{t,2}^{\text{swap}}) - y_{t,2}^{\text{nom}}$ and $W > 0$, a QE-driven fall in y^{TIPS} moves W toward zero (compresses $|W|$) even if π^{swap} and y^{nom} are unchanged. Second, rising 2-year breakeven inflation expectations ($y^{\text{nom}} - y^{\text{TIPS}}$) through early 2021 imply y^{nom} was rising relative to y^{TIPS} ; for W , if the breakeven rise exceeds the inflation swap rate adjustment, W falls—compressing $|W|$ from above given $W > 0$. Third, the 2-year tenor is among the thinnest TIPS maturities by outstanding notional and is more sensitive to idiosyncratic supply and demand factors than the 5- or 10-year tenors. The 5y and 10y TIPS series show near-zero and statistically insignificant exit responses, consistent with the interpretation that the TIPS-Treasury channel operates primarily through hedge fund rather than primary dealer balance sheets Siritwardane et al., 2025, Table VII, diluting the regulatory signal. The TIPS 2y anomaly does not overturn the main result in a statistical sense, but it substantially qualifies its interpretation. The pooled point estimate is driven almost entirely by the UST spot-futures subgroup: all three UST spot-futures series show significant re-widening, while the TIPS-Treasury series are mixed. A reader should therefore interpret $\hat{\beta} = +3.84$ bps as primarily measuring the UST spot-futures balance sheet channel, with the TIPS channel empirically unresolved in this sample window due to confounding from QE purchasing, rising breakeven inflation, and thin 2y TIPS issuance.

4.5 Robustness

Table 8 and Figure 6 present a systematic comparison of TOTAL and DIRECT specifications for both events, both window sizes, and both coefficients of interest. The key quantities in Table 8 are $\Delta\hat{\beta}$ and Δt . For the exit event, the Treasury interaction is highly stable across specifications. For the entry event, the Treasury interaction is similarly stable while the non-Treasury baseline shifts substantially as discussed in Section 4.2. This asymmetry is exactly what the balance sheet channel predicts: the SLR relief operates through leverage capacity, not through funding rates, so controlling for SOFR should affect the non-Treasury baseline without affecting the Treasury differential.

Spread series are winsorized at the first and 99th percentiles within each series prior to estimation. The winsorization log in Table 9 confirms symmetric clipping of approximately 2.0 to 2.1% of observations per series. Summing the clipped observations across all 17 series in the dependent variable $|W|$ yields approximately 1,396 clipped observations out of 68,620 (2.04 percent). The most consequential clip is the CIP-EUR minimum, which is bounded at -432.4 basis points (the March 2020 COVID spike). Without this bound, the extreme observation would exert disproportionate influence on the HAC covariance matrix during the entry window. All main results are quantitatively unchanged under alternative clipping

thresholds of p5/p95.

Newey-West HAC standard errors account for serial correlation within each series but assume cross-sectional independence. In my setting, all arbitrage spreads are exposed to common shocks (monetary policy, risk appetite), and residuals are plausibly correlated across series within the same trading day. I therefore re-estimate the pooled regressions using the Driscoll and Kraay, 1998 variance estimator, which is consistent under both

Table 8: Robustness: Total vs. Direct Specification Comparison

Event	Window	Coefficient	N	$\hat{\beta}_{TOT}$	SE_{TOT}	t_{TOT}	Stars	$\hat{\beta}_{DIR}$	SE_{DIR}	t_{DIR}	$\Delta\hat{\beta}$	Δt
Entry	$W = 20$	Post	697	-0.84	4.77	-0.18	—	+1.45	7.03	+0.21	+2.29	+0.38
Entry	$W = 20$	Post $\times$ Treas.	697	-2.16	6.13	-0.35	—	-2.23	6.14	-0.36	-0.06	-0.01
Entry	$W = 60$	Post	2,057	-7.08	3.05	-2.32	**	-2.76	5.54	-0.50	+4.32	+1.82
Entry	$W = 60$	Post $\times$ Treas.	2,057	-4.54	3.00	-1.51	—	-4.52	3.00	-1.51	+0.02	+0.00
Exit	$W = 20$	Post	697	-2.62	0.77	-3.40	***	-2.28	0.68	-3.33	+0.34	+0.06
Exit	$W = 20$	Post $\times$ Treas.	697	+3.91	0.92	+4.26	***	+3.90	0.92	+4.24	-0.01	-0.02
Exit	$W = 60$	Post	2,057	-3.09	0.83	-3.72	***	-3.23	0.85	-3.82	-0.14	-0.10
Exit	$W = 60$	Post $\times$ Treas.	2,057	+3.84	0.89	+4.32	***	+3.84	0.88	+4.34	0.00	+0.02

Notes. Each row reports both TOTAL and DIRECT estimates of the stated coefficient from regression (1). “Post $\times$ Treas.” abbreviates $Post_t \times TreasuryBased_t$. Stars correspond to the TOTAL specification; DIRECT significance is identical for all exit rows. HAC SE (Newey-West, 5 lags). *** $p < 0.01$, ** $p < 0.05$.

Table 9: Winsorization Log: Per-Series Clipping Counts

Series	Variable	p_1	bound	p_{99}	bound	$n_{clipped, lo}$	$n_{clipped, hi}$	% clipped
CIP-AUD	$ W $		0.25		44.60	40	40	2.04
CIP-CAD	$ W $		0.24		45.14	40	40	2.04
CIP-CHF	$ W $		0.91		103.72	40	40	2.05
CIP-EUR	$ W $		1.58		432.41	40	40	2.04
CIP-GBP	$ W $		0.81		55.50	40	40	2.04
CIP-JPY	$ W $		12.67		110.43	40	40	2.05
CIP-NZD	$ W $		0.52		41.22	40	40	2.04
CIP-SEK	$ W $		1.11		92.02	40	40	2.04
Eq. INDU	$ W $		3.07		126.51	38	38	2.03
Eq. NDX	$ W $		1.88		109.21	38	38	2.04
Eq. SPX	$ W $		3.53		111.44	38	38	2.04
TIPS 10y	$ W $		1.66		38.38	38	38	2.03
TIPS 2y	$ W $		0.45		42.05	38	38	2.04
TIPS 5y	$ W $		0.77		38.45	38	38	2.03
UST SF 10y	$ W $		0.51		227.33	50	50	2.01
UST SF 2y	$ W $		0.42		178.09	50	50	2.01
UST SF 5y	$ W $		0.27		179.19	50	50	2.01

Notes. Each series is winsorized within-series at the first and 99th percentiles prior to all estimation. p_1 and p_{99} are the clipping bounds in basis points. “% clipped” reports the fraction of observations replaced at either tail. Total clipped observations in the $|W|$ dependent variable across the full stacked panel of 68,620 series-day pairs: approximately 1,396 (2.04 percent). The pipeline also winsorizes the signed spread W separately; counting both variables gives approximately 2,792 total clips, but the relevant figure for the estimation sample is 1,396.

temporal autocorrelation and unrestricted cross-sectional dependence. Standard errors are computed using the Bartlett kernel with five lags, matching the main specification. The finding that DK standard errors for the exit Treasury interaction are smaller than Newey-West standard errors (0.68 vs. 0.89 bps) deserves explanation. Driscoll-Kraay SEs are often described as more conservative than panel NW SEs, but that comparison applies when NW is computed treating each series observation as independent (understating cross-sectional dependence). In this stacked-panel setup, the NW estimator is applied to the full pooled sample of $N = 2,057$ series-day observations, treating each as an independent




Figure 6: Forest plot of pooled difference-in-differences estimates. Each row reports the $\text{Post} \times \text{TreasuryBased}$ interaction coefficient $\hat{\beta}$ with 95% HAC confidence interval (Newey-West, 5 lags) for a given event, window, and control specification. The top panel covers the exit event (March 31, 2021); the bottom panel the entry event (April 1, 2020). The DiD estimate ($\hat{\mu}_2$, relief period) from Table 6 is included for reference. All exit estimates are positive and significant; all entry estimates are negative but insignificant.

draw over time. This overstates the effective time dimension by a factor of 17, producing upward-biased NW SEs relative to the true time-only precision. The DK estimator, by first averaging score contributions cross-sectionally, treats each calendar day as a single effective observation. For the $\text{Post} \times \text{TreasuryBased}$ interaction, which is identified by the differential post-event movement of the two groups, this cross-sectional averaging removes common-factor noise while retaining the between-group differential that identifies $\hat{\beta}$. At entry, the opposite pattern holds: the COVID crash induced strong co-movement across all arbitrage series, which DK correctly penalizes (DK SE = 6.32 vs. NW SE = 3.00 bps), correctly reducing the t -statistic from -1.51 to -0.72 .

Table 10 reports DK standard errors for the key coefficients at $W = 60$. For the exit event, the Treasury interaction yields $t_{\text{DK}} = +5.62$ (TOTAL) and $+5.65$ (DIRECT), compared to $t_{\text{NW}} = +4.32$ and $+4.34$ under Newey-West. The DK standard error is smaller than the NW standard error (0.683 vs. 0.889 bps) because the cross-sectional averaging of score contributions removes the common-factor noise while retaining the series-specific identifying variation. For the entry event, the DK standard error for the Treasury interaction is substantially larger (6.32 vs. 3.00 bps), reducing t_{DK} to -0.72 , consistent with the interpretation that the COVID crash induced strong cross-sectional co-movement that NW does not fully correct for. The non-Treasury baseline at entry remains significant under DK for the TOTAL specification ($t_{\text{DK}} = -2.62$), but not for DIRECT ($t_{\text{DK}} = -0.64$), consistent with the interpretation that SOFR controls absorb the funding-channel co-movement that drives the entry-period compression of non-Treasury spreads.

Table 10: Driscoll-Kraay Robustness: $W = 60$ Pooled Regressions

Event	Spec	Coefficient	Coef.	SE _{NW}	t_{NW}	SE _{DK}	t_{DK}
<i>Exit event (2021-03-31)</i>							
Exit	TOTAL	Post	-3.09	0.830	-3.72***	0.957	-3.22***
Exit	TOTAL	Post $\times$ Treasury	+3.84	0.889	+4.32***	0.683	+5.62***
Exit	DIRECT	Post	-3.23	0.846	-3.82***	0.912	-3.54***
Exit	DIRECT	Post $\times$ Treasury	+3.84	0.885	+4.34***	0.680	+5.65***
<i>Entry event (2020-04-01)</i>							
Entry	TOTAL	Post	-7.08	3.054	-2.32**	2.699	-2.62***
Entry	TOTAL	Post $\times$ Treasury	-4.54	3.003	-1.51	6.323	-0.72
Entry	DIRECT	Post	-2.76	5.541	-0.50	4.331	-0.64
Entry	DIRECT	Post $\times$ Treasury	-4.52	2.995	-1.51	6.315	-0.72

Notes. All regressions use $N = 2,057$ series-day observations (17 series $\times$ 121 trading-day window). SE_{NW}: Newey-West HAC (5 lags). SE_{DK}: Driscoll and Kraay, 1998 with Bartlett kernel, 5 lags, computed via the score-summing approach allowing unrestricted cross-sectional dependence. The exit Treasury interaction remains significant at the one-percent level under both estimators; the DK t -statistic exceeds the NW t -statistic because cross-sectional averaging of score contributions reduces common-factor noise. *** $p < 0.01$, ** $p < 0.05$.

5 Discussion

The identification limitations of this approach warrant explicit discussion. First, the SLR relief is a single nationwide policy change, not a staggered rollout. The Callaway and Sant’Anna, 2021 estimator is designed for settings with variation in treatment timing across units, and its motivating concern—contaminated comparisons in two-way fixed effects with staggered adoption—does not apply here. My design is closer to the MMF reform event study in Siriwardane et al., 2025, Section 4.2 and the classical single-treatment DiD of the program evaluation literature. The relevant identification threat is instead whether some other event contemporaneous with the SLR entry or exit differentially affected Treasury-based versus non-Treasury spreads. The parallel-trends evidence from the DiD 2019 baseline (Section 4.1) provides the sharpest rebuttal to this concern: Treasury-based and non-Treasury spreads were on the same trend before the SLR event and diverged only when the relief took effect. Second, as detailed in Section 4.2, the pre-period for the entry event coincides with the COVID-19 stress episode. I address this by treating the exit event as the primary source of evidence and the clean DiD (Section 4.1) as the whole-sample complement. Third, the CIP-EUR series contains a single extreme observation at 432.4 basis points in March 2020, driven by the near-breakdown of the cross-currency basis market. This observation is winsorized at the within-series 99th-percentile bound of 432.41 bps (see Table 9). My exit-event estimates use none of this extreme observation in their sample window. Fourth, the DIRECT specification includes rolling 7-, 14-, and 30-calendar-day total Treasury issuance constructed from auction records by maturity bucket. These controls are included to isolate the balance sheet channel from mechanical Treasury supply effects. Their estimated coefficients are economically small and statistically indistinguishable from zero across all exit specifications, and inclusion of these controls does not meaningfully change the key Treasury interaction (see Table 8). The absence of a supply effect on the Treasury differential is consistent with the balance sheet channel: the SLR relief operates through intermediary leverage capacity, not through the supply of Treasury collateral. Fifth, Cochran et al., 2023 examine dealer Treasury holdings from regulatory reports and find no noticeable change in aggregate dealer Treasury inventory attributable to the SLR exclusion, which they interpret as evidence against a supply-constraint channel. My results are not necessarily inconsistent with theirs: aggregate inventory is a different object from spread levels, and a marginal reduction in dealer balance sheet costs can tighten arbitrage spreads without requiring a discrete change in reported inventory positions. Moreover, my identification relies on the differential response of Treasury versus non-Treasury spreads at two specific event dates, which is a sharper test than the aggregate inventory comparison. I acknowledge this as an empirical tension worth investigating with more granular Y-9C

or FR 2004 position data. A sixth technical concern is that five Newey-West lags may be insufficient given AR(1) coefficients exceeding 0.93, at which lag-5 autocorrelation is approximately 0.70. Applying the rule of thumb of $T^{1/4}$ lags for a sample of 121 trading days suggests 3–4 lags for the event-window regressions, consistent with the choice of five. The Driscoll-Kraay results (which strengthen the exit Treasury interaction to $t_{DK} = 5.62$) provide reassurance that the primary inference is not sensitive to the autocorrelation correction, though a formal sensitivity analysis with 20 or more lags would further address this concern for the long time-series panels used in the DiD.

The results shown in this paper invite a macroeconomic reading that the paper’s design supports only in a specific manner, and it is worth being explicit about which conclusions the evidence provides and which it does not. To be clear, the SLR exclusion is a single nationwide regulatory event that the 2019 pre-trend test does not endorse as a clean natural experiment (Section 4); concurrent quantitative easing, emergency repo operations, and the run-up to the Standing Repo Facility all overlap with the relief window. What the pattern of results does establish is narrower but, I argue, still informative for several active debates in macro-finance. The first concerns Treasury market functioning as monetary-policy infrastructure. Since the September 2019 repo spike and the March 2020 dash-for-cash, the FOMC and the Group of Thirty have treated dealer balance-sheet capacity as a first-order constraint on the Treasury market’s ability to absorb intermediation flows (Liang & Parkinson, 2020). Two competing readings of these episodes have circulated: dealer constraints are essentially a stress phenomenon that binds during liquidity crises but is slack in normal times, and the constraint is a continuously priced friction whose shadow cost shows up in arbitrage spreads even outside acute episodes. The differential compression of Treasury spreads during the SLR relief window is hard to reconcile with the former. If the leverage-ratio constraint were slack in normal times, removing it should not have moved spreads materially during the calmer second half of the relief window; the data show that it did. The estimates are consistent with the continuous-friction reading and contribute a price-based benchmark that is absent from the existing literature, which has relied primarily on quantity-based evidence from dealer inventories and Y-9C balance-sheet data. The second debate concerns the design of post-2008 capital regulation. The non-risk-weighted leverage ratio was adopted as a backstop to the risk-weighted framework, deliberately chosen to be insensitive to asset composition so that banks could not optimize their reported capital ratios by tilting toward low-risk-weight assets. An unintended consequence, noted from the rule’s earliest critiques, is that dealers face the same denominator charge for U.S. Treasuries as for unsecured corporate loans (Favara et al., 2022; Liang & Parkinson, 2020). Whether this charge is economically material has been debated on theoretical grounds; the SLR exclusion provides one of the few opportunities to read the magnitude directly off market

prices. The 70% compression of UST spot-futures spreads during relief, and the +3.84 basis point re-widening at expiry, suggest that the implicit tax that the leverage ratio levies on Treasury intermediation is not negligible at the daily-pricing margin. This finding has direct bearing on the Federal Reserve’s 2025 proposal to recalibrate the enhanced SLR buffer for G-SIBs: the 2020–2021 episode is the only quasi-experimental observation of the price-impact magnitude such a reform should be expected to deliver, and the estimates here suggest the magnitude is several basis points on Treasury-based arbitrage spreads on a sustained basis.

The most direct contribution of this paper is to the segmented arbitrage research program of Siriwardane et al., 2025. They identify balance-sheet segmentation through two events: the 2012 London Whale episode, in which JPMorgan’s risk-limit tightening differentially widened the equity spot-futures spreads it intermediated as well as hedge-fund return correlations linking specific arbitrage spreads to specific intermediary types. The 2020–2021 SLR exclusion provides a third event study of the same logical structure, with three features that complement their evidence. First, it is the symmetric inverse of the London Whale shock: a positive balance-sheet shock that compresses spreads in the trades the targeted intermediary intermediates, then a reversal at expiry. Second, the targeted intermediary type (covered bank holding companies) and the targeted asset class (U.S. Treasuries and central bank reserves) are specified by the regulation itself, not inferred from data, eliminating the ambiguity present in the London Whale episode about which trades JPMorgan was actually marginal in. Third, the event falls outside the Siriwardane et al., 2025 sample (which ends February 2020), so the result extends their identification logic into the COVID and post-COVID regulatory regime: a regime in which both quantitative easing and the post-Basel-III leverage ratio operate at full force. The pattern this paper documents is the predicted signature of balance-sheet segmentation in their framework, and it operates over an episode and a shock structure they did not study. The SLR event is, by construction, a balance-sheet-constraint shock rather than a funding-source shock: it removed Treasuries and reserves from the leverage denominator without altering the secured repo or unsecured funding markets through which the trades are financed. The DIRECT specification isolates the balance-sheet-constraint channel from contemporaneous secured-funding-rate movements. The Treasury interaction $\hat{\beta}$ is stable both at the exit event across TOTAL and DIRECT specifications and in the full-period DiD across the same specifications. This indicates that the differential Treasury response is orthogonal to secured-funding conditions. The reading consistent with the data is balance-sheet segmentation specifically, not funding segmentation.

Two caveats discipline these readings. First, the comparison is across strategy types, not across institutions: the paper does not estimate which fraction of the pooled -10.76

bps operates through covered bank holding companies’ SLR-binding status versus through general equilibrium spillovers (freed dealer capacity, hedge-fund leverage shifts, anticipatory portfolio rebalancing into Treasuries during the relief horizon). A bank-level mechanism test using FR Y-9C data, sketched in Section 6 as the principal extension, is the direct test of the channel attribution. Second, the entire exercise rests on a regulatory event whose effect cannot be separated from concurrent QE without out-of-sample variation. Replication around the 2025 enhanced-SLR recalibration, when implemented, would provide the second observation that this single-event design lacks.

What this paper offers is not an answer to the segmented arbitrage problem but a structured starting point for it. The SLR exclusion delivers a regulatory shock that is targeted by construction, and the cross-strategy panel delivers a falsification benchmark by construction; together, they let one read off a price-based magnitude for the leverage-ratio friction in Treasury arbitrage that is robust to the specific funding-rate alternative and to the specific cross-sectional dependence correction. That magnitude is consistent with the segmentation framework, broadly consistent with the dealer-constraint estimates of Bräuning and Stein, 2024, and at variance with the representative-intermediary view in its strongest form. Whether those readings will hold up under the bank-level mechanism test, the next regulatory episode, or both, is the question the paper leaves open.

6 Conclusion

This paper investigates whether regulatory leverage constraints are the operative friction behind Treasury market arbitrage dislocations. Using the Federal Reserve’s 2020–2021 SLR exclusion as a temporary regulatory event, I find that the expiry of regulatory relief is associated with Treasury-based arbitrage spreads widening by 3.84 basis points more than non-Treasury spreads in the subsequent 60 trading days. A full-period difference-in-differences estimates that Treasury spreads compressed by 10.76 basis points more than non-Treasury spreads during the relief window, relative to the 2019 pre-COVID baseline. Non-Treasury series show no comparable break at the SLR event dates, passing the falsification tests that are the design’s primary identification safeguard.

These results extend the segmented arbitrage research program of Siriwardane et al., 2025 along three dimensions. The authors identify balance-sheet segmentation in part through the 2012 London Whale episode, a single negative balance-sheet shock to one intermediary. The 2020–2021 SLR exclusion is a positive balance-sheet shock to a regulatorily-defined class of intermediaries, operates over a one-year window with both onset and expiry observable, and falls entirely outside the authors’ February 2020 sample cutoff. The pattern of differential compression is the signature balance-sheet segmentation predicts under their

framework, identified on a shock and a horizon their evidence did not include.

The results also speak to an active policy debate. In 2025, the Federal Reserve proposed reforms to the enhanced SLR buffer for G-SIBs that would free substantial capital at G-SIB depository subsidiaries. The 2020–2021 episode provides the only available quasi-experimental benchmark for the price-impact magnitude of such a reform. My estimates suggest that a comparable structural change in leverage capacity could compress Treasury-based arbitrage spreads by several basis points on a sustained basis, reducing the intermediation cost embedded in Treasury market prices.

Three extensions remain for future work. First, a bank-level mechanism test using quarterly FR Y-9C data would directly link the aggregate spread compression to variation in individual banks’ SLR bind, providing the cross-sectional identification that the current design cannot. Second, the point estimate from the full-period DiD, referred to as “partial reversion” in Section 4.1, implies approximately 16% of the relief-period compression reversed in the nine-month post-expiry window, though this figure is statistically imprecise. This suggests that the structural impact of the exclusion, combined with concurrent Fed balance sheet expansion and the launch of the Standing Repo Facility, may have durably improved Treasury market intermediation in ways that outlasted the formal regulation. Third, a local projections specification would allow the impulse response of each spread to be estimated flexibly over post-event horizons, which would be informative about the speed of arbitrage capital reallocation following the regulatory change.

Data and Code Availability Statement

The replication package for this paper is publicly available at https://github.com/BaileyMeche/slr_bucket. The repository contains the full data construction pipeline, all econometric code, and the scripts that generate every table and figure reported in the paper. The original sources from which these cleaned panels are constructed are documented in Table 1; end-to-end reproduction of the data can be done from the listed sources using the pipeline in the repository.

References

- Adrian, T., Etula, E., & Muir, T. (2014). Financial intermediaries and the cross-section of asset returns. *Journal of Finance*, 69(6), 2557–2596.
- Andersen, L., Duffie, D., & Song, Y. (2019). Funding value adjustments. *Journal of Finance*, 74(1), 145–192.
- Bräuning, F., & Stein, H. (2024). *The effect of primary dealer constraints on intermediation in the Treasury market* (Research Department Working Paper No. 24-7). Federal Reserve Bank of Boston.
- Brunnermeier, M. K., & Pedersen, L. H. (2009). Market liquidity and funding liquidity. *The Review of Financial Studies*, 22(6), 2201–2238.
- Callaway, B., & Sant’Anna, P. H. C. (2021). Difference-in-differences with multiple time periods. *Journal of Econometrics*, 225(2), 200–230.
- Cochran, P., Infante, S., Petrasek, L., Saravay, Z., & Tian, M. H. (2023). *Dealers’ Treasury market intermediation and the supplementary leverage ratio* (FEDS Notes). Board of Governors of the Federal Reserve System.
- Driscoll, J. C., & Kraay, A. C. (1998). Consistent covariance matrix estimation with spatially dependent panel data. *Review of Economics and Statistics*, 80(4), 549–560.
- Du, W., Tepper, A., & Verdelhan, A. (2018). Deviations from covered interest rate parity. *Journal of Finance*, 73(3), 915–957.
- Favara, G., Infante, S., & Rezende, M. (2022). *Leverage regulations and Treasury market participation: Evidence from credit line drawdowns* [Working paper, available at SSRN 4175429.].
- Fleckenstein, M., & Longstaff, F. A. (2020). Renting balance sheet space: Intermediary balance sheet rental costs and the valuation of derivatives. *The Review of Financial Studies*, 33(11), 5051–5091.
- Fleckenstein, M., Longstaff, F. A., & Lustig, H. (2014). The TIPS-Treasury bond puzzle. *Journal of Finance*, 69(5), 2151–2197.
- Gârleanu, N., & Pedersen, L. H. (2011). Margin-based asset pricing and deviations from the law of one price. *The Review of Financial Studies*, 24(6), 1980–2022.
- Gromb, D., & Vayanos, D. (2002). Equilibrium and welfare in markets with financially constrained arbitrageurs. *Journal of Financial Economics*, 66(2–3), 361–407.
- He, Z., Kelly, B., & Manela, A. (2017). Intermediary asset pricing: New evidence from many asset classes. *Journal of Financial Economics*, 126(1), 1–35.
- He, Z., & Krishnamurthy, A. (2013). Intermediary asset pricing. *American Economic Review*, 103(2), 732–770.

- Liang, N., & Parkinson, P. (2020). *Enhancing liquidity of the U.S. Treasury market under stress* (Hutchins Center Working Paper No. 72). Brookings Institution, Hutchins Center on Fiscal and Monetary Policy.
- Shleifer, A., & Vishny, R. W. (1997). The limits of arbitrage. *The Journal of Finance*, 52(1), 35–55.
- Siriwardane, E. N., Sunderam, A., & Wallen, J. (2025). Segmented arbitrage. *Journal of Finance*, 80(5), 2543–2590.